

SALARY PREDICTION USING MACHINE LEARNING


PRESENTATION BY:

AKHILESH MAURYA



PROBLEM STATEMENTS

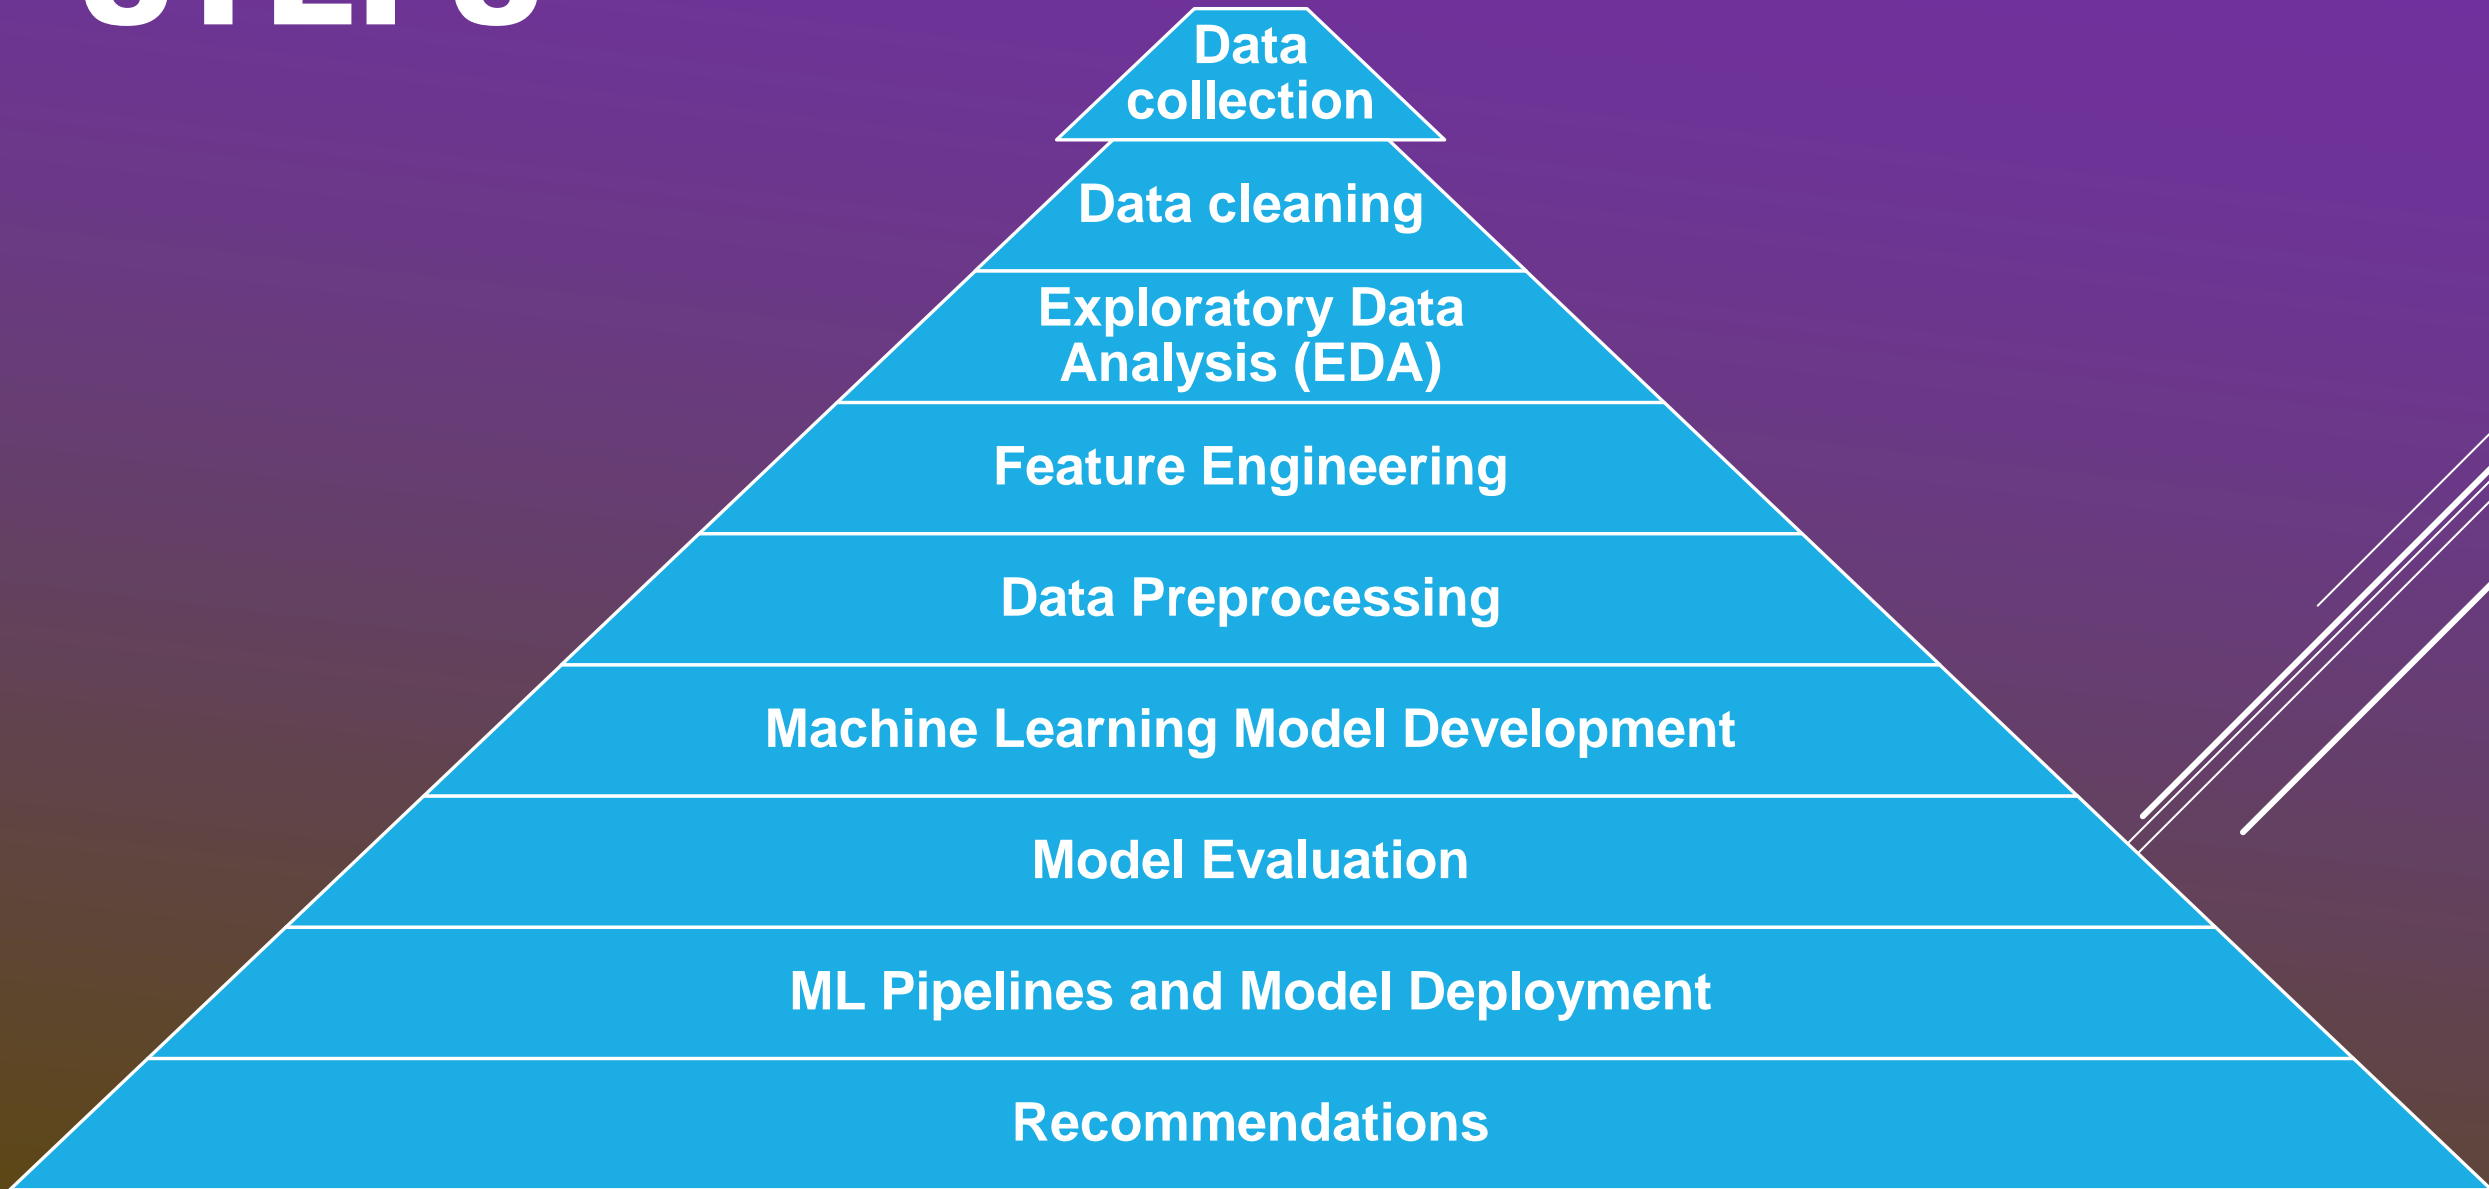
Salaries in the field of data professions vary widely based on factors such as experience, job role, and performance. Accurately predicting salaries for data professionals is essential for both job seekers and employers.

Several thin, white, parallel diagonal lines are positioned in the bottom right corner of the slide, extending from the right edge towards the center.

DATASET OVERVIEW

- ☐ **`FIRST NAME`**: First name
- ☐ **`LAST NAME`**: Last name
- ☐ **`SEX`**: Gender
- ☐ **`DOJ`**: Date of joining the company
- ☐ **`CURRENT DATE`**: Current date of data
- ☐ **`DESIGNATION`**: Job role/designation
- ☐ **`AGE`**: Age
- ☐ **`SALARY`**: Target variable, the salary of the data professional
- ☐ **`UNIT`**: Business unit or department
- ☐ **`LEAVES USED`**: Number of leaves used
- ☐ **`LEAVES REMAINING`**: Number of leaves remaining
- ☐ **`RATINGS`**: Ratings or performance ratings
- ☐ **`PAST EXP`**: Past work experience

STEPS



1. DATA COLLECTION

+ Code + Text

 df.head()



	FIRST NAME	LAST NAME	SEX	DOJ	CURRENT DATE	DESIGNATION	AGE	SALARY	UNIT	LEAVES USED	LEAVES REMAINING	RATINGS	PAST EXP
0	TOMASA	ARMEN	F	5-18-2014	01-07-2016	Analyst	21.0	44570	Finance	24.0	6.0	2.0	0
1	ANNIE	NaN	F	NaN	01-07-2016	Associate	NaN	89207	Web	NaN	13.0	NaN	7
2	OLIVE	ANCY	F	7-28-2014	01-07-2016	Analyst	21.0	40955	Finance	23.0	7.0	3.0	0
3	CHERRY	AQUILAR	F	04-03-2013	01-07-2016	Analyst	22.0	45550	IT	22.0	8.0	3.0	0
4	LEON	ABOULAHOU	M	11-20-2014	01-07-2016	Analyst	NaN	43161	Operations	27.0	3.0	NaN	3

2. DATA CLEANING

✓ Convert Data Types

```
▶ df['DOJ'] = pd.to_datetime(df['DOJ'], errors='coerce')
df['CURRENT DATE'] = pd.to_datetime(df['CURRENT DATE'], errors='coerce')
df['SEX'] = df['SEX'].astype('category')
df['DESIGNATION'] = df['DESIGNATION'].astype('category')
df['UNIT'] = df['UNIT'].astype('category')
```


✓ filling the missing values

```
[ ] df['LAST NAME'].fillna(method='ffill', inplace=True)
df['DOJ'].fillna(method='bfill', inplace=True)
df['AGE'].fillna(df['AGE'].median(), inplace=True)
df['LEAVES USED'].fillna(df['LEAVES USED'].median(), inplace=True)
df['LEAVES REMAINING'].fillna(df['LEAVES REMAINING'].median(), inplace=True)
df['RATINGS'].fillna(df['RATINGS'].median(), inplace=True)
```

3. EXPLORATORY DATA ANALYSIS

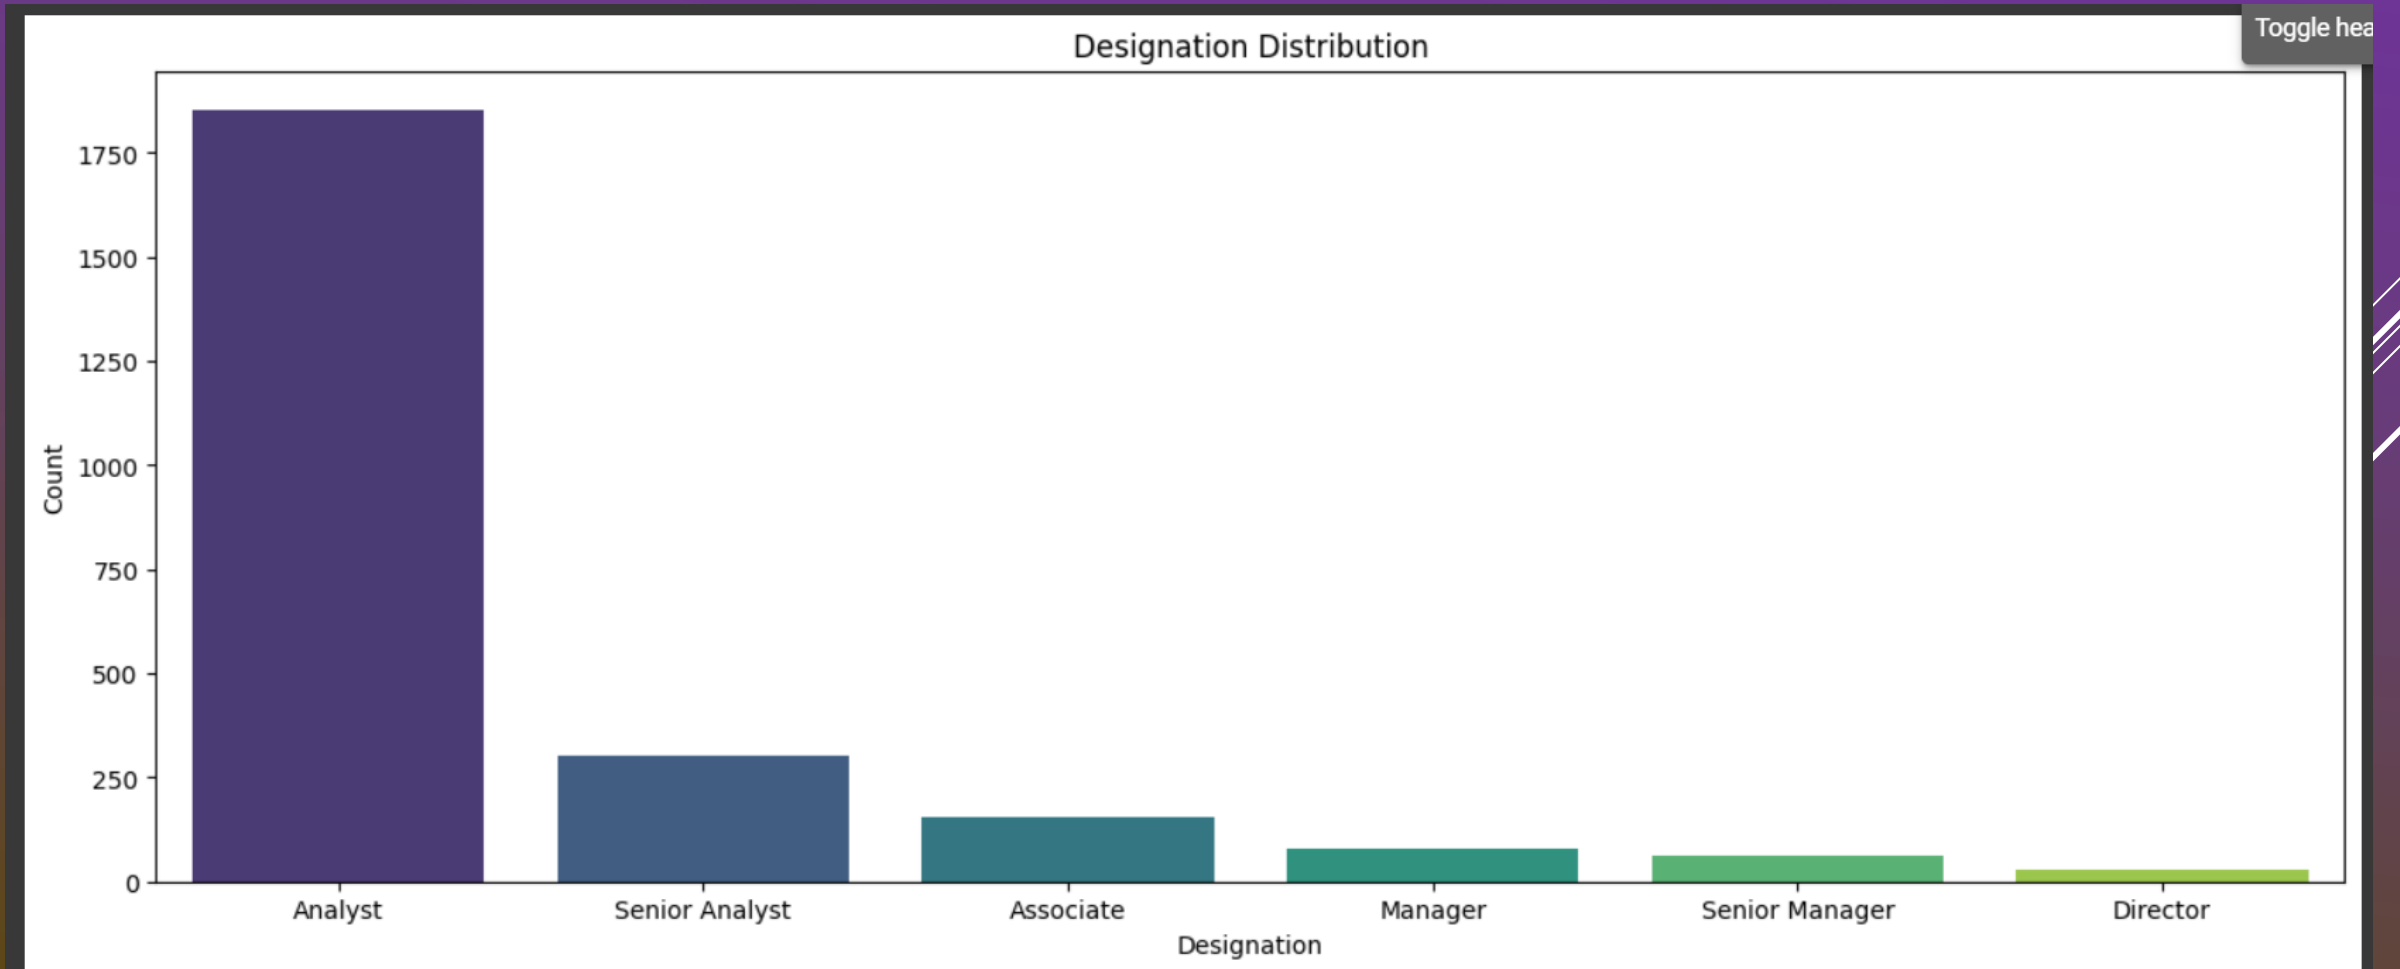
► Gender

```
] df['SEX'].value_counts()  
plt.figure(figsize=(8, 6))  
plt.pie(df['SEX'].value_counts(), labels=df['SEX'].value_counts().index, autopct='%1.1f%%')  
plt.show()
```



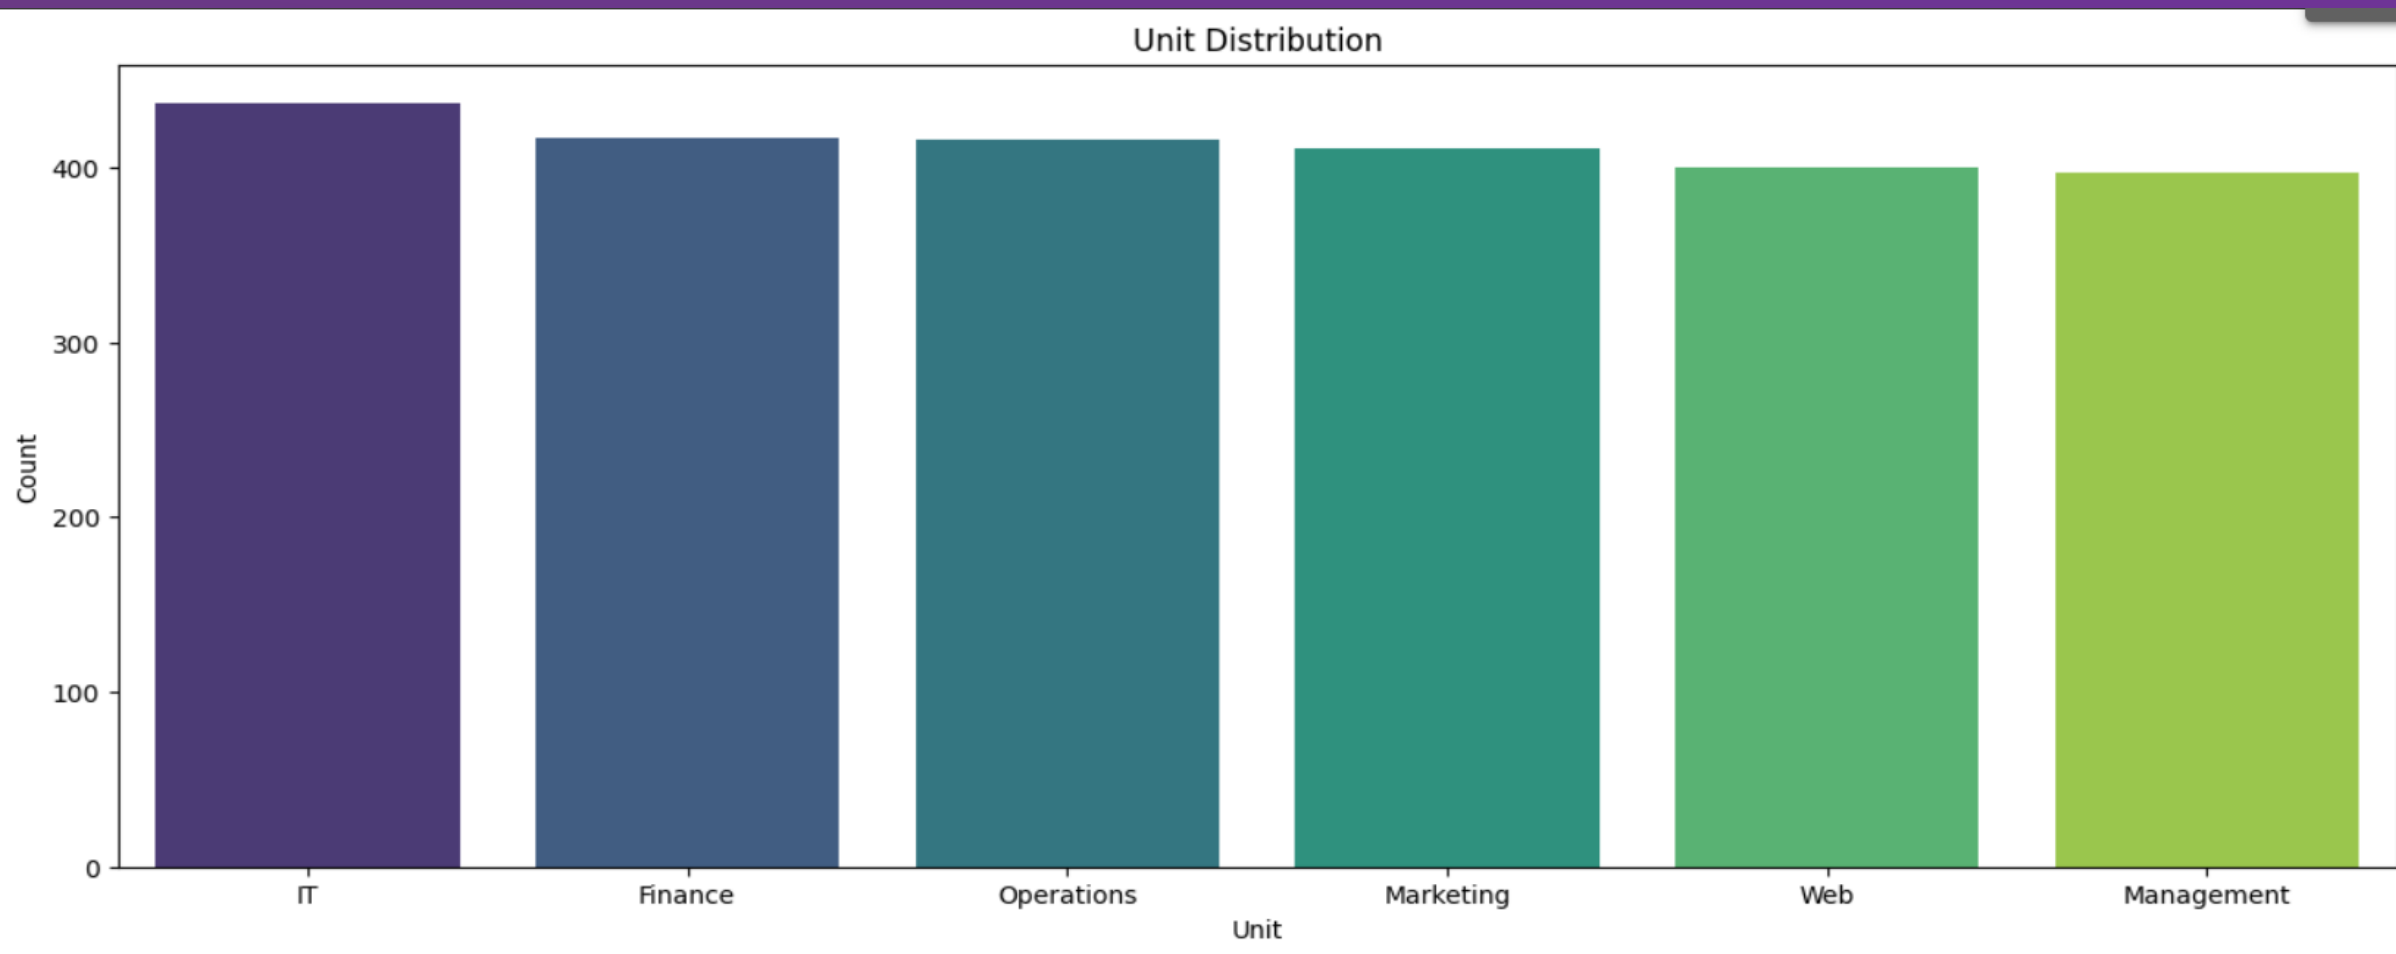
EXPLORATORY DATA ANALYSIS

► Designation



EXPLORATORY DATA ANALYSIS


► Unit



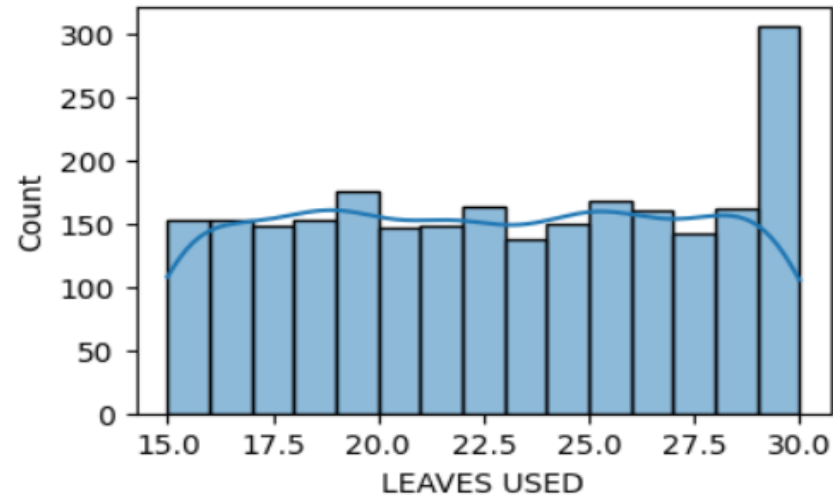
EXPLORATORY DATA ANALYSIS

Numerical columns Distribution

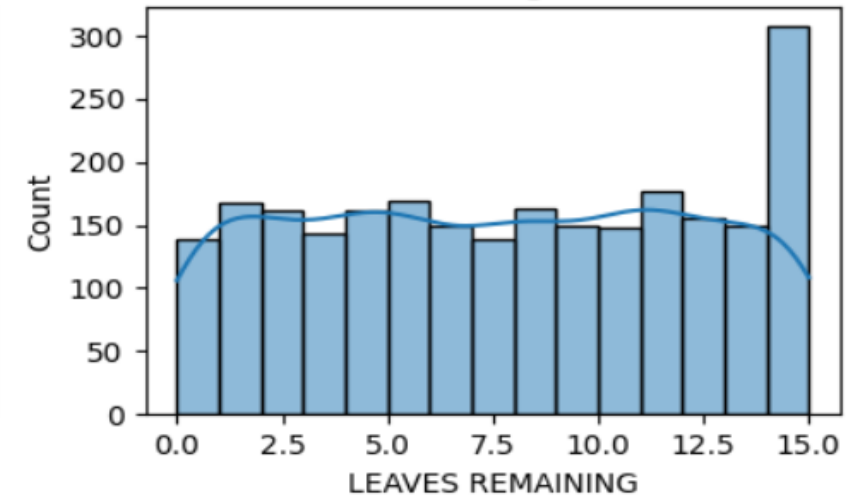
Age Distribution



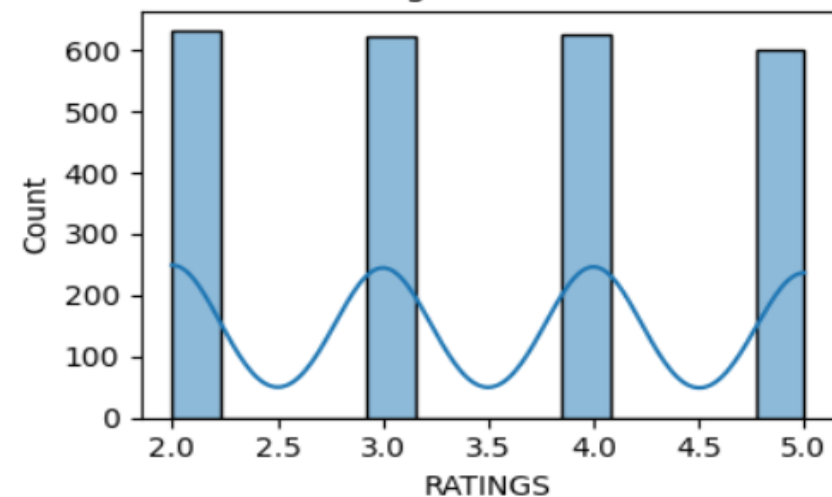
Leaves Used Distribution




Leaves Remaining Distribution



Ratings Distribution




Salary Distribution




EXPLORATORY DATA ANALYSIS

Outliers


Age Distribution




Leaves Used Distribution




Leaves Remaining Distribution



Ratings Distribution




Salary Distribution




EXPLORATORY DATA ANALYSIS

Outliers Removing


Age Distribution




Leaves Used Distribution




Leaves Remaining Distribution



Ratings Distribution



Salary Distribution



4. FEATURE ENGINEERING

Creating new columns

4. Feature Engineering

```
[ ] df['TOTAL LEAVES'] = df['LEAVES USED'] + df['LEAVES REMAINING']
```

```
[ ] df['EXPERIENCE'] = (df['CURRENT DATE'] - df['DOJ']).dt.days // 365
```

```
[ ] df.head(2)
```



	FIRST NAME	LAST NAME	SEX	DOJ	CURRENT DATE	DESIGNATION	AGE	SALARY	UNIT	LEAVES USED	LEAVES REMAINING	RATINGS	PAST EXP	TOTAL LEAVES	EXPERIENCE
0	TOMASA	ARMEN	F	2014-05-18	2016-01-07	Analyst	21.0	44570.0	Finance	24.0	6.0	2.0	0	30.0	1
1	ANNIE	ARMEN	F	2014-07-28	2016-01-07	Associate	24.0	60707.5	Web	22.0	13.0	3.0	7	35.0	1

4. DATA PREPROCESSING

Creating Pipeline

```
[ ] from sklearn.pipeline import Pipeline
    from sklearn.preprocessing import StandardScaler, OneHotEncoder
    from sklearn.compose import ColumnTransformer
    from sklearn.linear_model import LinearRegression
    from sklearn.tree import DecisionTreeRegressor
    from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
    from sklearn.svm import SVR
    from xgboost import XGBRegressor
    from sklearn.model_selection import GridSearchCV
    from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
    from sklearn.model_selection import train_test_split
```

✓ Prepare Data for Modeling

```
[ ] categorical_features = ['SEX', 'DESIGNATION', 'UNIT']
    numeric_features = ['AGE', 'RATINGS', 'PAST EXP', 'TOTAL LEAVES', 'EXPERIENCE']

    x = df.drop('SALARY', axis=1)
    y = df['SALARY']

    x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

5. MACHINE LEARNING MODEL

Creating Pipeline for Multiple model

```
preprocessor = ColumnTransformer(  
    transformers=[  
        ('num', StandardScaler(), numeric_features),  
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features)  
    ]  
)  
  
# Define the pipelines  
pipelines = {  
    'Linear Regression': Pipeline(steps=[  
        ('preprocessor', preprocessor),  
        ('model', LinearRegression())  
    ]),  
    'Decision Tree': Pipeline(steps=[  
        ('preprocessor', preprocessor),  
        ('model', DecisionTreeRegressor(random_state=42))  
    ]),  
    'Random Forest': Pipeline(steps=[  
        ('preprocessor', preprocessor),  
        ('model', RandomForestRegressor(random_state=42))  
    ]),  
    'Gradient Boosting': Pipeline(steps=[  
        ('preprocessor', preprocessor),  
        ('model', GradientBoostingRegressor(random_state=42))  
    ]),  
    'SVR': Pipeline(steps=[  
        ('preprocessor', preprocessor),  
        ('model', SVR())  
    ]),  
    'XGBoost': Pipeline(steps=[  
        ('preprocessor', preprocessor),  
        ('model', XGBRegressor(random_state=42))  
    ])  
}
```

6. MODEL EVALUATION

Checking mae, mse, r2-score

```
# Evaluate each model
results = {}
for model_name, pipeline in pipelines.items():
    pipeline.fit(x_train, y_train)
    y_pred = pipeline.predict(x_test)
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    results[model_name] = {'MAE': mae, 'MSE': mse, 'R^2': r2}

# Print the results
for model_name, metrics in results.items():
    print(f"{model_name} - MAE: {metrics['MAE']}, MSE: {metrics['MSE']}, R^2: {metrics['R^2']}")
```

```
Linear Regression - MAE: 2330.5875225769655, MSE: 8216047.127319162, R^2: 0.8314757580314052
Decision Tree - MAE: 2727.993754800307, MSE: 13498296.06087705, R^2: 0.7231284002786358
Random Forest - MAE: 2531.766590639536, MSE: 10540254.673026312, R^2: 0.7838025511049753
Gradient Boosting - MAE: 2350.764210874021, MSE: 8410040.49263974, R^2: 0.8274966444344443
SVR - MAE: 5386.135914852255, MSE: 51006566.37711205, R^2: -0.04622609886713103
XGBoost - MAE: 2563.1857516381046, MSE: 10677913.116747579, R^2: 0.7809789566782144
```


6. HYPERPARAMETER TUNING

Find best model accuracy

```
best_models = {}
for model_name, pipeline in pipelines.items():
    grid_search = GridSearchCV(pipeline, param_grids[model_name], cv=5, scoring='neg_mean_squared_error', n_jobs=-1)
    grid_search.fit(X_train, y_train)
    best_models[model_name] = grid_search.best_estimator_
    print(f"{model_name} Best Params: {grid_search.best_params_}")
```

```
Linear Regression Best Params: {'model__fit_intercept': True}
Decision Tree Best Params: {'model__max_depth': 5, 'model__min_samples_split': 20}
Random Forest Best Params: {'model__max_depth': 5, 'model__min_samples_split': 10, 'model__n_estimators': 100}
Gradient Boosting Best Params: {'model__learning_rate': 0.05, 'model__max_depth': 3, 'model__n_estimators': 100}
SVR Best Params: {'model__C': 10, 'model__epsilon': 0.2, 'model__kernel': 'linear'}
XGBoost Best Params: {'model__learning_rate': 0.05, 'model__max_depth': 3, 'model__n_estimators': 100}
```


```
# Evaluate the tuned models
tuned_results = {}
for model_name, model in best_models.items():
    y_pred = model.predict(X_test)
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)
    tuned_results[model_name] = {'MAE': mae, 'MSE': mse, 'R^2': r2}
```

```
for model_name, metrics in tuned_results.items():
    print(f"{model_name} - MAE: {metrics['MAE']}, MSE: {metrics['MSE']}, R^2: {metrics['R^2']}")
```

```
Linear Regression - MAE: 2330.5875225769655, MSE: 8216047.127319162, R^2: 0.8314757580314052
Decision Tree - MAE: 2314.3054221126145, MSE: 8401563.662765643, R^2: 0.8276705177468403
Random Forest - MAE: 2312.2913113925924, MSE: 8258674.385095295, R^2: 0.8306014049279479
Gradient Boosting - MAE: 2343.492062354923, MSE: 8362709.723221706, R^2: 0.8284674740699562
SVR - MAE: 3144.7309812089134, MSE: 14877534.722720956, R^2: 0.6948380136268658
XGBoost - MAE: 2343.234926285282, MSE: 8362059.5587196965, R^2: 0.8284808099817589
```

6. HYPERPARAMETER TUNING

Here can see best model in graph



7. FINAL RECOMMENDATIONS

Here can see best model in code

```
def predict_salary(input_data):  
    # Load the saved model  
    model = joblib.load('best_model.pkl')  
  
    # Preprocess the input data in the same way as the training data  
    return model.predict(input_data)
```

```
# Example input data for prediction  
example_input_data = pd.DataFrame([  
    'SEX': 'M',  
    'DESIGNATION': 'Analyst',  
    'AGE': 25,  
    'UNIT': 'Finance',  
    'RATINGS': 4,  
    'PAST EXP': 3,  
    'TOTAL LEAVES': 30,  
    'EXPERIENCE': 3  
])
```

```
predicted_salary = predict_salary(example_input_data)  
predicted_salary
```

```
array([44957.74319876])
```

8. CONCLUSION

- Linear Regression is the best-performing model with an R^2 score of 0.8315, indicating it explains approximately 83% of the variance in salary.
- Random Forest and Gradient Boosting models also performed well with R^2 scores of 0.8306 and 0.8285 respectively.
- SVR had the lowest performance with an R^2 score of 0.6948, suggesting it is less effective for this task.
- Overall, Linear Regression is recommended for salary prediction due to its strong performance across key metrics.